

IBM Advance Data Science Capstone Project

Import Required Libraries

In [19]:

```
# import data wrangling and visual libraries
import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv), data manipulation as in SQL
import matplotlib.pyplot as plt # this is used for the plot the graph
import seaborn as sns # used for plot interactive graph. I like it most for plot
%matplotlib inline

# to check the dirctory
import os
print(os.listdir("../Data"))
```

Loading Data

In [67]:

```
# Read Field Names File and Show Top 2 Records
with open('../Data/field_names.txt', 'r') as the_file:
    col_names = [line.strip() for line in the_file.readlines()]
print(col_names, '\n')

# Read Data File and Show Top 2 Records
data = pd.read_csv("../Data/breast-cancer.csv", header=None)
print(data.head(2))

['ID', 'diagnosis', 'radius_mean', 'radius_sd_error', 'radius_worst', 'texture_mean', 'texture_sd_error', 'texture_worst', 'perimeter_mean', 'perimeter_sd_error', 'perimeter_worst', 'area_mean', 'area_sd_error', 'area_worst', 'smoothness_mean', 'smoothness_sd_error', 'smoothness_worst', 'compactness_mean', 'compactness_sd_error', 'compactness_worst', 'concavity_mean', 'concavity_sd_error', 'concavity_worst', 'concave_points_mean', 'concave_points_sd_error', 'concave_points_worst', 'symmetry_mean', 'symmetry_sd_error', 'symmetry_worst', 'fractal_dimension_mean', 'fractal_dimension_sd_error', 'fractal_dimension_worst']
```

	0	1	2	3	4	5	6	7	8	9	\
0	842302	M	17.99	10.38	122.8	1001.0	0.11840	0.27760	0.3001	0.14710	
1	842517	M	20.57	17.77	132.9	1326.0	0.08474	0.07864	0.0869	0.07017	
	...		22	23	24	25	26	27	28	29	\
0	...		25.38	17.33	184.6	2019.0	0.1622	0.6656	0.7119	0.2654	
1	...		24.99	23.41	158.8	1956.0	0.1238	0.1866	0.2416	0.1860	
	30	31									
0	0.4601	0.11890									
1	0.2750	0.08902									

[2 rows x 32 columns]

In [68]:

```
# set Column Names and display top 3 records
```

```
data.columns = col_names
print(data.head(3))
# to see last 3 records
print(data.tail(3))
```

```

      ID diagnosis  radius_mean  radius_sd_error  radius_worst \
0      842302      M      17.99      10.38      122.8
1      842517      M      20.57      17.77      132.9
2  84300903      M      19.69      21.25      130.0

      texture_mean  texture_sd_error  texture_worst  perimeter_mean \
0      1001.0      0.11840      0.27760      0.3001
1      1326.0      0.08474      0.07864      0.0869
2      1203.0      0.10960      0.15990      0.1974

      perimeter_sd_error      ...      concavity_worst \
0      0.14710      ...      25.38
1      0.07017      ...      24.99
2      0.12790      ...      23.57

      concave_points_mean  concave_points_sd_error  concave_points_worst \
0      17.33      184.6      2019.0
1      23.41      158.8      1956.0
2      25.53      152.5      1709.0

      symmetry_mean  symmetry_sd_error  symmetry_worst  fractal_dimension_mean\
0      0.1622      0.6656      0.7119  0.2654
1      0.1238      0.1866      0.2416      0.1860
2      0.1444      0.4245      0.4504      0.2430

      fractal_dimension_sd_error  fractal_dimension_worst
0      0.4601      0.11890
1      0.2750      0.08902
2      0.3613      0.08758
```

[3 rows x 32 columns]

```

      ID diagnosis  radius_mean  radius_sd_error  radius_worst \
566  926954      M      16.60      28.08      108.30
567  927241      M      20.60      29.33      140.10
568  92751      B      7.76      24.54      47.92

      texture_mean  texture_sd_error  texture_worst  perimeter_mean \
566      858.1      0.08455      0.10230      0.09251
567      1265.0      0.11780      0.27700      0.35140
568      181.0      0.05263      0.04362      0.00000

      perimeter_sd_error      ...      concavity_worst \
566      0.05302      ...      18.980
567      0.15200      ...      25.740
568      0.00000      ...      9.456

      concave_points_mean  concave_points_sd_error  concave_points_worst \
566      34.12      126.70      1124.0
567      39.42      184.60      1821.0
568      30.37      59.16      268.6

      symmetry_mean  symmetry_sd_error  symmetry_worst  fractal_dimension_mean \
566      0.11390      0.30940      0.3403      0.1418
567      0.16500      0.86810      0.9387      0.2650
568      0.08996      0.06444      0.0000      0.0000

      fractal_dimension_sd_error  fractal_dimension_worst
566      0.2218      0.07820
567      0.4087      0.12400
568      0.2871      0.07039
```

[3 rows x 32 columns]

Data Wrangling

We have successfully loaded data. Now lets look at the type of data we have.

In [52]:

```
data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 32 columns):
ID                                     569 non-null int64
diagnosis                             569 non-null object
radius_mean                           569 non-null float64
radius_sd_error                       569 non-null float64
radius_worst                          569 non-null float64
texture_mean                          569 non-null float64
texture_sd_error                      569 non-null float64
texture_worst                         569 non-null float64
perimeter_mean                       569 non-null float64
perimeter_sd_error                   569 non-null float64
perimeter_worst                     569 non-null float64
area_mean                            569 non-null float64
area_sd_error                        569 non-null float64
area_worst                           569 non-null float64
smoothness_mean                      569 non-null float64
smoothness_sd_error                  569 non-null float64
smoothness_worst                    569 non-null float64
compactness_mean                     569 non-null float64
compactness_sd_error                 569 non-null float64
compactness_worst                    569 non-null float64
concavity_mean                       569 non-null float64
concavity_sd_error                   569 non-null float64
concavity_worst                      569 non-null float64
concave_points_mean                  569 non-null float64
concave_points_sd_error              569 non-null float64
concave_points_worst                 569 non-null float64
symmetry_mean                        569 non-null float64
symmetry_sd_error                    569 non-null float64
symmetry_worst                       569 non-null float64
fractal_dimension_mean               569 non-null float64
fractal_dimension_sd_error           569 non-null float64
fractal_dimension_worst              569 non-null float64
dtypes: float64(30), int64(1), object(1)
memory usage: 142.3+ KB
```

So we have 569 records against 32 columns and all of them have 569 non-null records and the data type is float64.

Lets keep the Diagnosis data and drop ID and diagnosis columns as they are not needed

In [57]:

```
# y includes our labels and x includes our features
y = data.diagnosis # M or B
list = ['ID', 'diagnosis']
x = data.drop(list,axis = 1 )
x.head()
```

Remaining features are representing the 3 computations (Mean, Standard Deviation Error and Worst) against single feautre. Lets group them into 3 categories.

In [84]:

```
#list of column names that match with Mean
mean_cols = [col for col in x.columns if '_mean' in col]
print(mean_cols, '\n')

#list of column names that match with SD
sd_cols = [col for col in x.columns if '_sd' in col]
print(mean_cols, '\n')
```

```
#list of column names that match with Mean
worst_cols = [col for col in x.columns if '_worst' in col]
print(mean_cols)
```

```
['radius_mean', 'texture_mean', 'perimeter_mean', 'area_mean', 'smoothness_mean', 'compactness_mean', 'concavity_mean', 'concave_points_mean', 'symmetry_mean', 'fractal_dimension_mean']
```

```
['radius_mean', 'texture_mean', 'perimeter_mean', 'area_mean', 'smoothness_mean', 'compactness_mean', 'concavity_mean', 'concave_points_mean', 'symmetry_mean', 'fractal_dimension_mean']
```

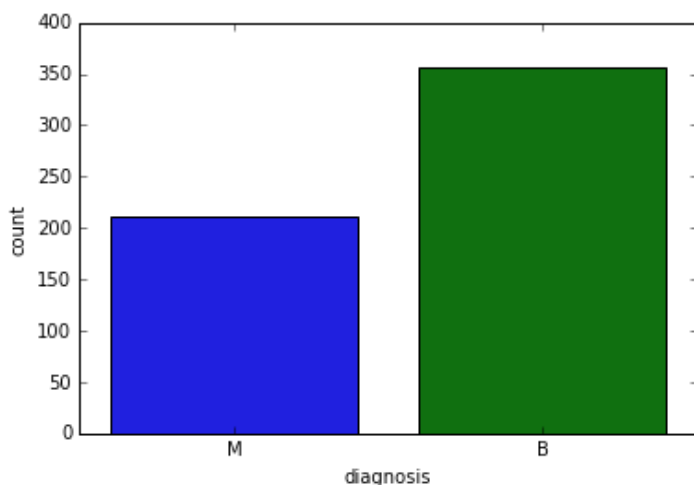
```
['radius_mean', 'texture_mean', 'perimeter_mean', 'area_mean', 'smoothness_mean', 'compactness_mean', 'concavity_mean', 'concave_points_mean', 'symmetry_mean', 'fractal_dimension_mean']
```

Before moving forward to the data analysis of features. Lets have a quick look of the labels we are going to predict

In [71]:

```
ax = sns.countplot(y,label="Count") # M = 212, B = 357
B, M = y.value_counts()
print('Number of Benign: ',B)
print('Number of Malignant : ',M)
```

Number of Benign: 357
Number of Malignant : 212



So far, we are unfamiliar with the data and its features, and what they are representing. In real world, we come across to many different problems where we don't know the meaning of features but to imagine in our minds. What we must know is the distribution of data like **variance, standard deviation, number of sample (count) or max min values**. These type of information helps to understand the data, how normally distributed it is, or it has skewed distribution.

In [85]:

```
x.describe()
```

Out[85]:

	radius_mean	radius_sd_error	radius_worst	texture_mean	texture_sd_error	texture_worst	perimeter_mean	perimeter_
count	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	56
mean	14.127292	19.289649	91.969033	654.889104	0.096360	0.104341	0.088799	
std	3.524049	4.301036	24.298981	351.914129	0.014064	0.052813	0.079720	
min	6.981000	9.710000	43.790000	143.500000	0.052630	0.019380	0.000000	
25%	11.700000	16.170000	75.170000	420.300000	0.086370	0.064920	0.029560	
50%	13.370000	18.840000	86.240000	551.100000	0.095870	0.092630	0.061540	

75%	radius_mean	radius_sd_error	radius_worst	texture_mean	texture_sd_error	texture_worst	perimeter_mean	perimeter_sd_error
max	28.110000	39.280000	188.500000	2501.000000	0.163400	0.345400	0.426800	

8 rows × 30 columns

The summary statistics helps us to understand is we need standirdization or normalization before visualization, feature selection or classificaiton.

Before moving to Exploratory Analysis, lets develop a function to generate bootstrap samples.

In [161]:

```
def bootstrap_resample(X, n=None):
    """ Bootstrap resample an array_like
    Parameters
    -----
    X : array_like
        data to resample
    n : int, optional
        length of resampled array, equal to len(X) if n==None
    Results
    -----
    returns X_resamples
    """
    if isinstance(X, pd.Series):
        X = X.copy()
        X.index = range(len(X.index))
    if n == None:
        n = len(X)

    resample_i = np.floor(np.random.rand(n)*len(X)).astype(int)
    X_resample = np.array(X[resample_i])
    return X_resample
```

In [181]:

```
# Create new df variable for resampled data
df_resampled = pd.DataFrame(index=df.index, columns=df.columns, dtype=df.dtypes)
for col in x.columns:
    df_resampled[col] = bootstrap_resample(x[col])

# original data
x.ix[:50,:50]
```

Out[181]:

	radius_mean	radius_sd_error	radius_worst	texture_mean	texture_sd_error	texture_worst	perimeter_mean	perimeter_sd_
0	17.990	10.38	122.80	1001.0	0.11840	0.27760	0.30010	0.14
1	20.570	17.77	132.90	1326.0	0.08474	0.07864	0.08690	0.07
2	19.690	21.25	130.00	1203.0	0.10960	0.15990	0.19740	0.12
3	11.420	20.38	77.58	386.1	0.14250	0.28390	0.24140	0.10
4	20.290	14.34	135.10	1297.0	0.10030	0.13280	0.19800	0.10
5	12.450	15.70	82.57	477.1	0.12780	0.17000	0.15780	0.08
6	18.250	19.98	119.60	1040.0	0.09463	0.10900	0.11270	0.07
7	13.710	20.83	90.20	577.9	0.11890	0.16450	0.09366	0.05
8	13.000	21.82	87.50	519.8	0.12730	0.19320	0.18590	0.09
9	12.460	24.04	83.97	475.9	0.11860	0.23960	0.22730	0.08
10	16.020	23.24	102.70	797.8	0.08206	0.06669	0.03299	0.03
11	15.780	17.89	103.60	781.0	0.09710	0.12920	0.09954	0.06
12	19.170	24.80	132.40	1123.0	0.09740	0.24580	0.20650	0.11

13	15.850	23.95	103.70	782.7	0.08401	0.10020	0.09938	0.0
	radius_mean	radius_sd_error	radius_worst	texture_mean	texture_sd_error	texture_worst	perimeter_mean	perimeter_sd
14	13.730	22.61	93.60	578.3	0.11310	0.22930	0.21280	0.08
15	14.540	27.54	96.73	658.8	0.11390	0.15950	0.16390	0.07
16	14.680	20.13	94.74	684.5	0.09867	0.07200	0.07395	0.05
17	16.130	20.68	108.10	798.8	0.11700	0.20220	0.17220	0.10
18	19.810	22.15	130.00	1260.0	0.09831	0.10270	0.14790	0.09
19	13.540	14.36	87.46	566.3	0.09779	0.08129	0.06664	0.04
20	13.080	15.71	85.63	520.0	0.10750	0.12700	0.04568	0.03
21	9.504	12.44	60.34	273.9	0.10240	0.06492	0.02956	0.02
22	15.340	14.26	102.50	704.4	0.10730	0.21350	0.20770	0.09
23	21.160	23.04	137.20	1404.0	0.09428	0.10220	0.10970	0.08
24	16.650	21.38	110.00	904.6	0.11210	0.14570	0.15250	0.09
25	17.140	16.40	116.00	912.7	0.11860	0.22760	0.22290	0.14
26	14.580	21.53	97.41	644.8	0.10540	0.18680	0.14250	0.08
27	18.610	20.25	122.10	1094.0	0.09440	0.10660	0.14900	0.07
28	15.300	25.27	102.40	732.4	0.10820	0.16970	0.16830	0.08
29	17.570	15.05	115.00	955.1	0.09847	0.11570	0.09875	0.07
30	18.630	25.11	124.80	1088.0	0.10640	0.18870	0.23190	0.12
31	11.840	18.70	77.93	440.6	0.11090	0.15160	0.12180	0.05
32	17.020	23.98	112.80	899.3	0.11970	0.14960	0.24170	0.12
33	19.270	26.47	127.90	1162.0	0.09401	0.17190	0.16570	0.07
34	16.130	17.88	107.00	807.2	0.10400	0.15590	0.13540	0.07
35	16.740	21.59	110.10	869.5	0.09610	0.13360	0.13480	0.06
36	14.250	21.72	93.63	633.0	0.09823	0.10980	0.13190	0.05
37	13.030	18.42	82.61	523.8	0.08983	0.03766	0.02562	0.02
38	14.990	25.20	95.54	698.8	0.09387	0.05131	0.02398	0.02
39	13.480	20.82	88.40	559.2	0.10160	0.12550	0.10630	0.05
40	13.440	21.58	86.18	563.0	0.08162	0.06031	0.03110	0.02
41	10.950	21.35	71.90	371.1	0.12270	0.12180	0.10440	0.05
42	19.070	24.81	128.30	1104.0	0.09081	0.21900	0.21070	0.09
43	13.280	20.28	87.32	545.2	0.10410	0.14360	0.09847	0.06
44	13.170	21.81	85.42	531.5	0.09714	0.10470	0.08259	0.05
45	18.650	17.60	123.70	1076.0	0.10990	0.16860	0.19740	0.10
46	8.196	16.84	51.71	201.9	0.08600	0.05943	0.01588	0.00
47	13.170	18.66	85.98	534.6	0.11580	0.12310	0.12260	0.07
48	12.050	14.63	78.04	449.3	0.10310	0.09092	0.06592	0.02
49	13.490	22.30	86.91	561.0	0.08752	0.07698	0.04751	0.03
50	11.760	21.60	74.72	427.9	0.08637	0.04966	0.01657	0.01

51 rows × 30 columns



In [182]:

```
#sample data
df_resampled.ix[:50,:50]
```

Out[182]:

concavity_worst	concave_points_worst	radius_worst	radius_mean	radius_sd_error	texture_mean	texture_sd_error	textur
-----------------	----------------------	--------------	-------------	-----------------	--------------	------------------	--------

id	concave			convex			radius			texture				
	count	cavity_worst	points_worst	count	concave_worst	points_worst	radius_worst	radius_mean	radius_sd	error	texture_mean	texture_sd	error	texture
1	17.560	24.560	1479.0	143.70	12.320	10.89	420.3	0.11220						
2	12.360		331.6	97.03	11.930	15.79	408.2	0.07966						
3	16.510		471.4	130.70	9.567	17.68	378.2	0.09714						
4	13.030		567.7	87.32	13.660	18.59	324.9	0.10490						
5	17.380		973.1	125.50	16.030	16.83	1878.0	0.12370						
6	12.780		474.2	78.04	19.730	27.08	992.1	0.11860						
7	14.400		939.7	76.83	9.755	16.85	1214.0	0.08974						
8	14.990		614.9	103.70	19.160	39.28	221.2	0.08223						
9	25.700		2232.0	66.20	8.888	14.78	1311.0	0.09989						
10	11.060		1025.0	82.50	9.777	18.84	633.0	0.09905						
11	13.780		806.9	84.06	11.200	24.04	556.7	0.10640						
12	21.200		367.0	78.99	11.900	20.28	1482.0	0.09029						
13	13.330		268.6	98.92	28.110	20.66	685.9	0.07838						
14	12.840		527.2	70.15	19.690	12.83	321.4	0.06995						
15	12.410		2384.0	91.43	9.567	16.68	800.0	0.08675						
16	11.600		762.6	73.16	9.504	18.14	748.9	0.07557						
17	14.190		808.2	82.69	11.500	27.08	552.4	0.09427						
18	11.250		302.0	66.52	16.130	16.70	250.5	0.07941						
19	11.920		869.3	121.30	10.510	19.24	349.6	0.11410						
20	16.890		2403.0	69.28	19.190	22.44	678.1	0.12480						
21	15.770		1866.0	129.70	14.060	21.35	507.6	0.10280						
22	11.150		826.0	88.73	15.100	14.96	538.4	0.09198						
23	11.980		436.6	87.38	12.540	12.22	448.6	0.08915						
24	22.930		2384.0	70.67	16.170	14.23	666.0	0.11840						
25	19.590		1872.0	100.30	13.430	15.70	515.9	0.10630						
26	17.790		1349.0	72.17	17.290	16.02	420.3	0.07376						
27	19.770		745.5	95.77	12.050	23.77	427.3	0.09950						
28	25.580		674.7	81.35	12.560	15.62	541.6	0.11700						
29	12.680		395.4	61.06	13.340	24.98	537.3	0.07937						
30	12.470		328.1	92.41	15.280	18.90	471.3	0.08054						
31	11.930		1646.0	85.63	12.700	18.00	651.0	0.10020						
32	14.200		543.4	92.41	7.729	14.96	920.6	0.09579						
33	15.400		1025.0	130.70	21.560	20.86	514.3	0.09168						
34	15.790		1600.0	98.17	10.750	13.98	458.4	0.09037						
35	15.750		760.2	94.25	13.340	16.94	537.9	0.09831						
36	13.740		300.2	93.86	21.090	16.85	477.4	0.10890						
37	11.020		591.2	97.65	13.710	19.34	857.6	0.08206						
38	12.760		1261.0	124.40	19.810	25.13	747.2	0.08801						
39	9.965		310.1	82.61	14.920	11.89	493.8	0.10070						
40	17.870		470.9	107.00	23.090	24.99	311.7	0.09879						
41	13.600		374.4	96.12	12.060	19.59	1052.0	0.10420						
42	36.040		856.9	71.24	13.170	26.29	1206.0	0.10120						
43	20.270		455.7	61.93	20.340	20.18	476.3	0.08817						
44	24.330		1298.0	85.48	12.200	21.01	481.9	0.16340						

45	concavity_worst	concave_points_worst	radius_worst	radius_mean	radius_sd_error	texture_mean	texture_sd_error	textur
46	10.510	719.8	75.17	11.300	16.67	571.8	0.10380	
47	17.010	351.9	89.46	19.190	19.12	320.8	0.10420	
48	25.300	811.3	76.31	9.683	12.96	201.9	0.12570	
49	18.810	626.9	60.11	11.140	18.77	1265.0	0.12370	
50	20.580	605.5	105.10	17.080	21.68	386.3	0.09699	

51 rows × 30 columns



In [192]:

```
print("\n radius_mean of Original & Sampled dataset")
x.radius_mean.mean(), df_resampled.radius_mean.mean()

radius_mean of Original & Sampled dataset
```

Out[192]:

(14.127291739894563, 14.144697715289995)

Visualization

For data visualization, we are going to use seaborn plots. Violin and Swarm plots usually helps us to understand data easily.

First, we need to perform normalization/standirdization of data. Because differences between values of features are very high to observe inplot

In [95]:

```
data = x
data_nor= (data - data.mean()) / (data.std()) # standardization
```

In [118]:


```
sns.set(style="whitegrid", palette="muted")
data = pd.concat([y,data_nor[mean_cols]],axis=1) # concat data to form new
data = pd.melt(data,id_vars="diagnosis",
               var_name="features",
               value_name='value')
plt.figure(figsize=(10,10))
import time
tic = time.time()
sns.swarmplot(x="features", y="value", hue="diagnosis", data=data)

plt.xticks(rotation=90)
```

Out[118]:

(array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9]), <a list of 10 Text xticklabel objects>)







In [119]:

```
data = pd.concat([y,data_nor[sd_cols]],axis=1)
data = pd.melt(data,id_vars="diagnosis",
               var_name="features",
               value_name='value')
plt.figure(figsize=(10,10))
sns.swarmplot(x="features", y="value", hue="diagnosis", data=data)
plt.xticks(rotation=90)
```

Out[119]:

(array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9]), <a list of 10 Text xticklabel objects>)






In [122]:

```
data = pd.concat([y,data_nor[worst_cols]],axis=1)
data = pd.melt(data,id_vars="diagnosis",
               var_name="features",
               value_name='value')
plt.figure(figsize=(10,10))
sns.swarmplot(x="features", y="value", hue="diagnosis", data=data)
toc = time.time()
plt.xticks(rotation=90)
print("swarm plot time: ", toc-tic , " s")
```

swarm plot time: 132.68085765838623 s




3 Plots Explanation: In the above three graphs, we can see variance more clearly. Let me ask you a question, in these three plots which feature looks like more evident in terms of classification? In my opinion concavity_worst in last swarm plot malignant and benign looks separately, not totally but mostly. However, area_mean, area_sd_error and area_worst all three looks like malignant and benign are mixed so it is hard to classify by using this feature.

In [96]:

```
# All mean features
data_diag = y
data = pd.concat([y,data_nor[mean_cols]],axis=1)
data = pd.melt(data,id_vars="diagnosis",
               var_name="features",
               value_name='value')
plt.figure(figsize=(10,10))
sns.violinplot(x="features", y="value", hue="diagnosis", data=data,split=True, inner="quart")
plt.xticks(rotation=90)
```

Out[96]:

(array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9]), <a list of 10 Text xticklabel objects>)



Plot Explanation: In texture_mean, perimeter_mean, concave_points_mean and symetry_mean features, median


of the Malignant and Benign looks like separated so it can be good for classification. However, in area_mean feature, median of the Malignant and Benign does not looks like separated so it does not gives good information for classification.

In [97]:

```
data = pd.concat([y,data_nor[sd_cols]],axis=1)
data = pd.melt(data,id_vars="diagnosis",
               var_name="features",
               value_name='value')
plt.figure(figsize=(10,10))
sns.violinplot(x="features", y="value", hue="diagnosis", data=data,split=True, inner="quart")
plt.xticks(rotation=90)
```

Out[97]:

(array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9]), <a list of 10 Text xticklabel objects>)




In [98]:

```
data = pd.concat([y,data_nor[worst_cols]],axis=1)
data = pd.melt(data,id_vars="diagnosis",
               var_name="features",
               value_name='value')
plt.figure(figsize=(10,10))
```

```
sns.violinplot(x="features", y="value", hue="diagnosis", data=data, split=True, inner="quart")
plt.xticks(rotation=90)
```

Out[98]:

(array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9]), <a list of 10 Text xticklabel objects>)



Plot Explanation: Lets interpret one more thing about plot above, variable of concavity_worst and concave_points_worst looks similar but how can we decide whether they are correlated with each other or not. It can not always be true but, if the features are correlated with each other then we can drop one of them.

In order to compare two features deeper, lets use joint plot. In the joint plot below, it is really correlated. Pearsonr value is correlation value and 1 is the highest. Therefore, 0.82 is looks enough to say that they are correlated.


In [102]:

```
sns.jointplot(x.loc[:, 'texture_worst'], x.loc[:, 'symmetry_worst'], kind="regg", color="#ce1414")
```

Out[102]:

<seaborn.axisgrid.JointGrid at 0x13f43886550>





In last violin plot, concavity_worst, concave_points_worst, and radius_worst also looks similar. Let's plot the pair grid plot to see if they are correlated.

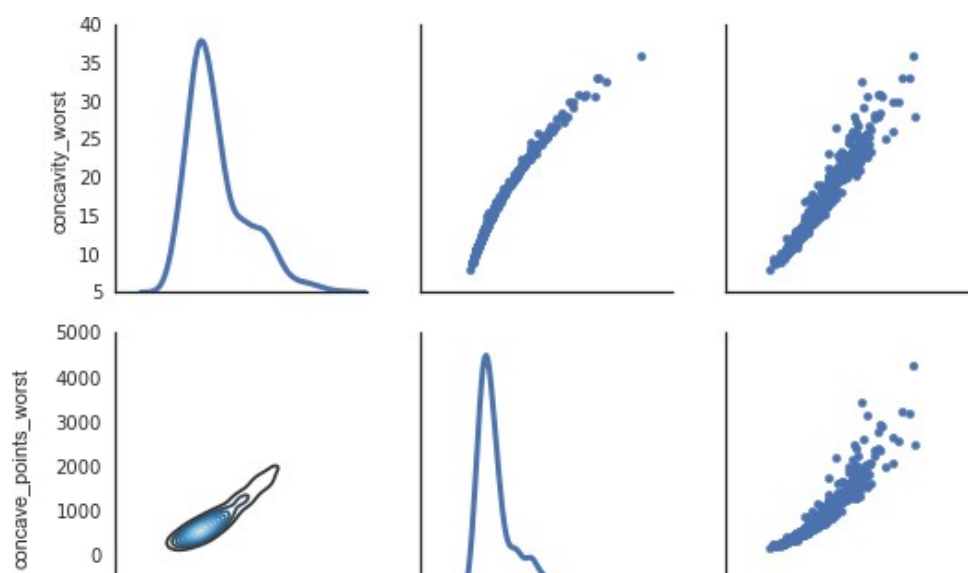
In [107]:

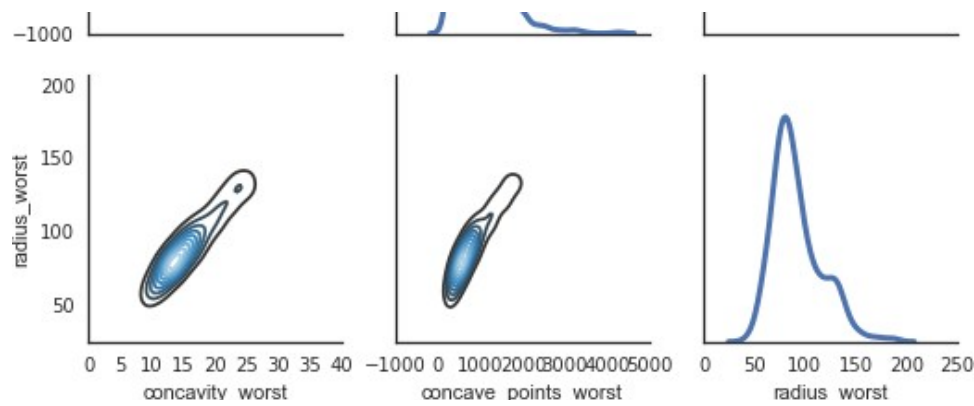
```
sns.set(style="white")
df = x.loc[:, ['concavity_worst', 'concave_points_worst', 'radius_worst']]
g = sns.PairGrid(df, diag_sharey=False)
g.map_lower(sns.kdeplot, cmap="Blues_d")
g.map_upper(plt.scatter)
g.map_diag(sns.kdeplot, lw=3)
```

```
C:\Users\sana.rasheed\AppData\Local\Continuum\Anaconda3\lib\site-packages\matplotlib\axes
\_axes.py:519: UserWarning: No labelled objects found. Use label='...' kwarg on individua
l plots.
  warnings.warn("No labelled objects found. ")
C:\Users\sana.rasheed\AppData\Local\Continuum\Anaconda3\lib\site-packages\matplotlib\axes
\_axes.py:519: UserWarning: No labelled objects found. Use label='...' kwarg on individua
l plots.
  warnings.warn("No labelled objects found. ")
C:\Users\sana.rasheed\AppData\Local\Continuum\Anaconda3\lib\site-packages\matplotlib\axes
\_axes.py:519: UserWarning: No labelled objects found. Use label='...' kwarg on individua
l plots.
  warnings.warn("No labelled objects found. ")
```

Out[107]:

<seaborn.axisgrid.PairGrid at 0x13f461d6048>





In above graph, we can observe that all three features are correlated.

Lets observe the correlation between all features and so we use this insights in feature selection for predictive model.

In [194]:

```
#correlation map
f,ax = plt.subplots(figsize=(18, 18))
sns.heatmap(x.corr(), annot=True, linewidths=.5, fmt= '.1f',ax=ax)
```

Out[194]:

<matplotlib.axes._subplots.AxesSubplot at 0x13f4a04b470>

To look into correlation matrix easily, next 3 plots are grouped by Mean, SD_Error and Worst

In [193]:

```
#correlation map
f,ax = plt.subplots(figsize=(18, 18))
sns.heatmap(x[mean_cols].corr(), annot=True, linewidths=.5, fmt= '.1f',ax=ax)
```

Out[193]:

<matplotlib.axes._subplots.AxesSubplot at 0x13f4a0380b8>


```
#correlation map
f,ax = plt.subplots(figsize=(18, 18))
sns.heatmap(x[sd_cols].corr(), annot=True, linewidths=.5, fmt= '.1f',ax=ax)
```

Out[195]:

<matplotlib.axes._subplots.AxesSubplot at 0x13f4930b898>

In [196]:

```
#correlation map
f,ax = plt.subplots(figsize=(18, 18))
sns.heatmap(x[worst_cols].corr(), annot=True, linewidths=.5, fmt= '.1f',ax=ax)
```

Out[196]:

<matplotlib.axes._subplots.AxesSubplot at 0x13f4bd83f60>

Feature Selection and Classification

In this Section, we will use 2 feature selection techniques and will test them with 2 Classifiers.

Following Feature Selection Techniques will be used:

1. Feature Selection by using correlation Heatmap
2. Univariate feature selection

In Classifications, we will use following 2 Techniques to train our model and predict:

1. Random Forest
2. Support Vector Machine

1. Feature Selection by using Correlation Heatmap

As it can be seen in map heat figure radius_mean, radius_worst, texture_mean, and concavity_worst are correlated with each others so we will use only concavity_worst. If you ask how I choose concavity_worst as a feature to use, well actually there is no correct answer, I just look at swarm plots and concavity_worst looks like clear for me. We cannot make exact separation among other correlated features without trying. So lets find other correlated features and look accuracy with random forest classifier.

The area_sd_error & smoothness_sd_error, perimeter_sd_error, concave_points_sd_error, concave_points_worst, texture_worst, symmetry_worst are few other highly correlated variables.

In [232]:

```
drop_list1 = [
    'radius_mean', 'radius_worst', 'texture_mean', 'concavity_worst',
    'area_sd_error', 'smoothness_sd_error', 'perimeter_sd_error',
    'concave_points_sd_error', 'concave_points_worst', 'texture_worst', 'symmetry_worst'
]
x_1 = x.drop(drop_list1,axis = 1 )          # do not modify x, we will use it later
x_1.head()
```

Out[232]:

	radius_sd_error	texture_sd_error	perimeter_mean	perimeter_worst	area_mean	area_worst	smoothness_mean	smoothnes
0	10.38	0.11840	0.3001	0.2419	0.07871	0.9053	8.589	0
1	17.77	0.08474	0.0869	0.1812	0.05667	0.7339	3.398	0
2	21.25	0.10960	0.1974	0.2069	0.05999	0.7869	4.585	0
3	20.38	0.14250	0.2414	0.2597	0.09744	1.1560	3.445	0
4	14.34	0.10030	0.1980	0.1809	0.05883	0.7813	5.438	0

Well, we choose our features but did we choose correctly ? Lets use random forest and find accuracy according to chosen features.

1.1 Random Forest Classification

In [257]:

```
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import f1_score, confusion_matrix
from sklearn.metrics import accuracy_score

# split data train 70 % and test 30%
x_train, x_test, y_train, y_test = train_test_split(x_1, y, test_size=0.3, random_state=
42)


#random forest classifier with n_estimators=10 (default)
clf_rf = RandomForestClassifier(random_state=43)
clr_rf = clf_rf.fit(x_train,y_train)

ac = accuracy_score(y_train,clf_rf.predict(x_train))
print('Random Forsest Accuracy is Training Data: ',ac)
cm = confusion_matrix(y_train,clf_rf.predict(x_train))
sns.heatmap(cm,annot=True,fmt="d")
```

Random Forsest Accuracy is Training Data: 0.989949748744

Out[257]:

<matplotlib.axes._subplots.AxesSubplot at 0x13f5039b2b0>






In [256]:

```
ac = accuracy_score(y_test,clf_rf.predict(x_test))
print('Random Forrest Accuracy on Test Data: ',ac)
cm = confusion_matrix(y_test,clf_rf.predict(x_test))
sns.heatmap(cm,annot=True,fmt="d")
```

Random Forrest Accuracy on Test Data: 0.93567251462

Out[256]:

<matplotlib.axes._subplots.AxesSubplot at 0x13f5021ad30>



The accuracy on Training Data is 98.9% and 93.5% on Test Data. We can see in confusion matrix that it has made few wrong predictions. Not bad, right!...

About the Overfit

Moreover, Accuracy gap between training and testing dataset is not wide, so **our model isn't overfitting**. If our model does much better on the training set than on the test set, then we're likely overfitting. For example, it would be a big red flag if our model saw 99% accuracy on the training set but only 55% accuracy on the test set.

There are few techniques which helps to prevent Overfitting:

1. Cross-validation: In standard k-fold cross-validation, we partition the data into k subsets, called folds. Then, we iteratively train the algorithm on k-1 folds while using the remaining fold as the test set (called the "holdout fold").
2. Train with more data: It won't work everytime, but training with more data can help algorithms detect the signal better. Of course, that's not always the case. If we just add more noisy data, this technique won't help.
3. Remove features: Some algorithms have built-in feature selection. For those that don't, we can manually improve their generalizability by removing irrelevant input features.
4. Early stopping: When we are training a learning algorithm iteratively, we can measure how well each iteration of the model performs. Up until a certain number of iterations, new iterations improve the model. After that point, however, the model's ability to generalize can weaken as it begins to overfit the training data. Early stopping refers stopping the training process before the learner passes that point.
5. Regularization: Regularization refers to a broad range of techniques for artificially forcing your model to be simpler. The method will depend on the type of learner you're using. For example, you could prune a decision tree, use dropout on a neural network, or add a penalty parameter to the cost function in regression. Oftentimes, the regularization method is a hyperparameter as well, which means it can be tuned through cross-validation.
6. Ensembling: Ensembles are machine learning methods for combining predictions from multiple separate models

Lets test these features with SVM classifier.

1.2 Support Vector Machine Classification

In [258]:

```
from sklearn import svm # for Support Vector Machine
from sklearn import metrics # for the check the error and accuracy of the model


model = svm.SVC()
model.fit(x_train,y_train)
#prediction=model.predict(x_test)
#metrics.accuracy_score(prediction,y_test)

ac = accuracy_score(y_test,model.predict(x_test))
print('Support Vector Machine Accuracy is: ',ac)
cm = confusion_matrix(y_test,model.predict(x_test))
sns.heatmap(cm,annot=True,fmt="d")
```

Support Vector Machine Accuracy is: 0.888888888889

Out[258]:

<matplotlib.axes._subplots.AxesSubplot at 0x13f503c6c50>



Accuracy is 88.8%, not good as compare to Random Forest model.

Now, lets test other feature selection methods if we could find better results.

By considering the correlation matrix we can select our desired feature set for the model, but machine learning domain is also equipped with few feature selection algorithms to extract/compute best feature set for our model. Here, we will test one of them:

2. Univariate Feature Selection

In univariate feature selection, we will use SelectKBest that removes all but the k highest scoring features.

http://scikit-learn.org/stable/modules/generated/sklearn.feature_selection.SelectKBest.html#sklearn.feature_selection.SelectKBest

In this method we need to choose how many features we will use. For example, will k (number of features) be 5 or 10 or 15? The answer is only try. We will not try all combinations but only choose k = 5 and find best 5 features for now.

In [259]:

```
from sklearn.feature_selection import SelectKBest
from sklearn.feature_selection import chi2

# split data train 70 % and test 30 %
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3, random_state=42)

# find best scored 5 features
```

```
select_feature = SelectKBest(chi2, k=5).fit(x_train, y_train)
print('Score list:', select_feature.scores_)
print('Feature list:', x_train.columns)
```

```
Score list: [ 1.77946492e+02  6.06916433e+01  1.34061092e+03  3.66899557e+04
 1.00015175e-01  3.41839493e+00  1.30547650e+01  7.09766457e+00
 1.95982847e-01  3.42575072e-04  2.45882967e+01  4.07131026e-02
 1.72696840e+02  6.12741067e+03  1.32470372e-03  3.74071521e-01
 6.92896719e-01  2.01587194e-01  1.39557806e-03  2.65927071e-03
 3.25782599e+02  1.16958562e+02  2.40512835e+03  7.50217341e+04
 2.63226314e-01  1.19077581e+01  2.58858117e+01  8.90751003e+00
 1.00635138e+00  1.23087347e-01]
```

```
Feature list: Index(['radius_mean', 'radius_sd_error', 'radius_worst', 'texture_mean',
 'texture_sd_error', 'texture_worst', 'perimeter_mean',
 'perimeter_sd_error', 'perimeter_worst', 'area_mean', 'area_sd_error',
 'area_worst', 'smoothness_mean', 'smoothness_sd_error',
 'smoothness_worst', 'compactness_mean', 'compactness_sd_error',
 'compactness_worst', 'concavity_mean', 'concavity_sd_error',
 'concavity_worst', 'concave_points_mean', 'concave_points_sd_error',
 'concave_points_worst', 'symmetry_mean', 'symmetry_sd_error',
 'symmetry_worst', 'fractal_dimension_mean',
 'fractal_dimension_sd_error', 'fractal_dimension_worst'],
 dtype='object')
```

The Top scored features are **fractal_dimension_mean**, **concave_points_worst**, **perimeter_sd_error**, **compactness_sd_error**, and **smoothness_sd_error**.

Let's build model on top of these 5 features.

2.1 Random Forest Classification


In [260]:

```
x_train_2 = select_feature.transform(x_train)
x_test_2 = select_feature.transform(x_test)
#random forest classifier with n_estimators=10 (default)
clf_rf_2 = RandomForestClassifier(n_estimators=100)
clr_rf_2 = clf_rf_2.fit(x_train_2,y_train)
ac_2 = accuracy_score(y_test,clf_rf_2.predict(x_test_2))
print('Random Forsest Accuracy is: ',ac_2)
cm_2 = confusion_matrix(y_test,clf_rf_2.predict(x_test_2))
sns.heatmap(cm_2,annot=True,fmt="d")
```

Random Forsest Accuracy is: 0.970760233918

Out[260]:

<matplotlib.axes._subplots.AxesSubplot at 0x13f4fc3fb70>



These top 5 features are play significant role in the data set. Accuracy is 97% and as it can be seen in confusion matrix, we have few wrong predictions. If we set `n_estimators=10` (default), then accuracy is 95% `n_estimators=50`, then accuracy is 96%

So, we also need to spend a good time just for model tuning .

2.2 Suport Vector Machine

In [261]:

```
from sklearn import svm # for Support Vector Machine
from sklearn import metrics # for the check the error and accuracy of the model

model = svm.SVC()
model.fit(x_train,y_train)

ac = accuracy_score(y_test,model.predict(x_test))
print('Support Vector Machine Accuracy is: ',ac)
cm = confusion_matrix(y_test,model.predict(x_test))
sns.heatmap(cm,annot=True,fmt="d")
```

Support Vector Machine Accuracy is: 0.631578947368

Out[261]:

<matplotlib.axes._subplots.AxesSubplot at 0x13f51588ba8>

Let's do the little model tuning to gain some good accuracey.

In [262]:

```
model = svm.SVC(kernel = 'linear',C=.1, gamma=10, probability = True)
model.fit(x_train,y_train)

ac = accuracy_score(y_test,model.predict(x_test))
print('Support Vector Machine Accuracy is: ',ac)
cm = confusion_matrix(y_test,model.predict(x_test))
sns.heatmap(cm,annot=True,fmt="d")
```

Support Vector Machine Accuracy is: 0.964912280702

